

# R Primer

Download at <https://github.com/gabrielrvsc/class>

# Getting Started with R

- ▶ You can download R from <https://www.r-project.org/>
- ▶ R is an open-source language mostly used for Statistics and data analysis.
- ▶ It is based on a core build with basic functions plus packages that can be installed.
- ▶ A package is a set of functions/data.
- ▶ R community is huge and you should use it! Most of your questions were already answered on Stack Overflow (<https://stackoverflow.com>).
- ▶ Stack Overflow is a forum where people post questions about programming.

# RStudio

- ▶ R alone is just a console that looks like a notepad. It is not very friendly.
- ▶ RStudio is an Integrated development environment (IDE) made for R.
- ▶ You can download RStudio from <https://rstudio.com/products/rstudio/download/>
- ▶ Make sure to get the free desktop version.
- ▶ RStudio will automatically find your R once installed.

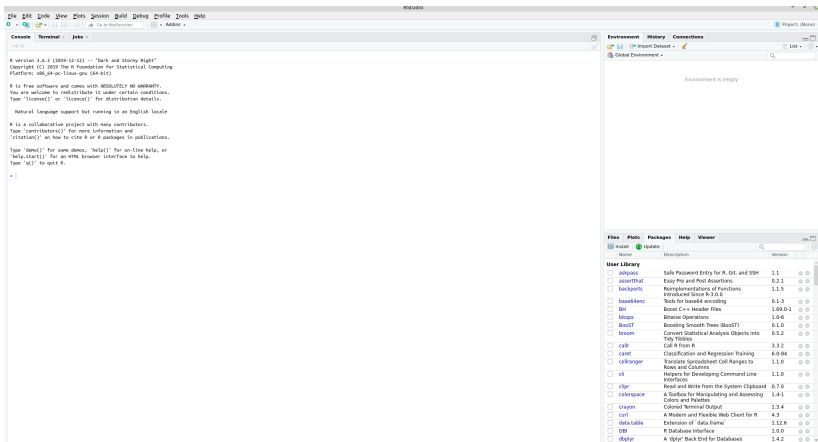
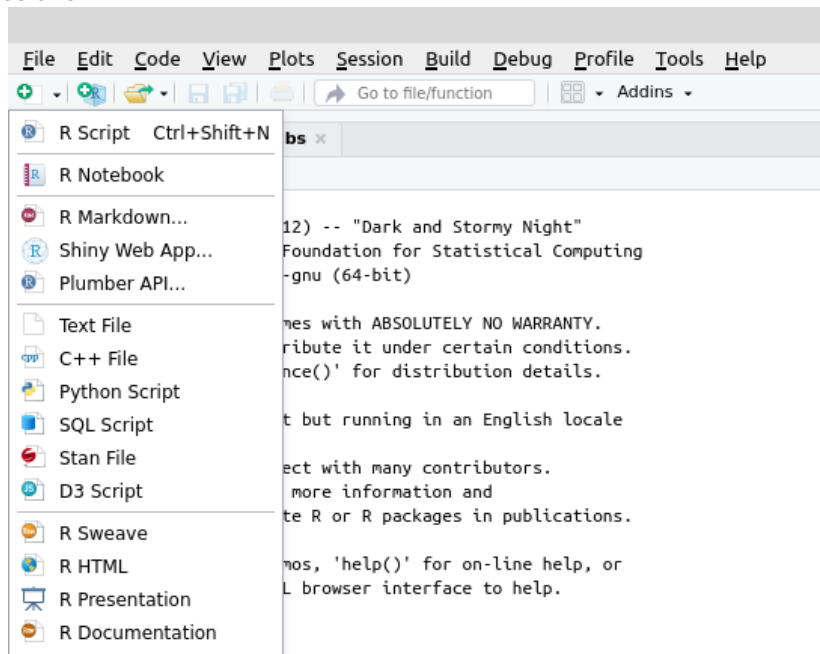


Figure 1: RStudio

# RStudio



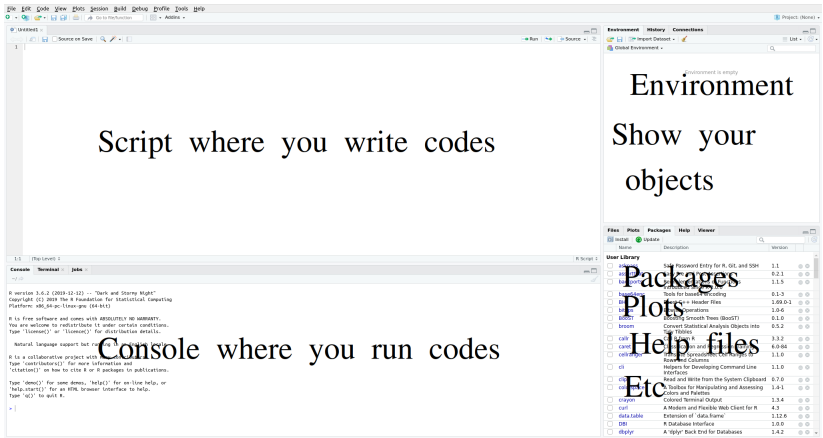


Figure 3: RStudio

# RStudio

- ▶ You can run commands on the console. For example, try `1+1` and hit enter.
- ▶ You can write codes on the script and run in the console by selecting them and pressing `Ctrl + Enter` (`Cmd+Enter` for MAC). You can also run the selected codes by pressing the Run button on the top right corner of the script.
- ▶ Type `1+1` on the script and run it.

# Creating Objects

- ▶ Objects are created with the symbols `=` or `<-`.
- ▶ creating an object is equivalent to storing the result.

```
x = 1
```

```
y <- 1
```

```
z = x + y
```



# Creating Objects

- The objects you create will show in your environment.

Global Environment ▾		<input type="text"/>
Values		
x	1	
y	1	
z	2	

# Objects

- ▶ You can see all your objects with:

```
ls()
```

```
## [1] "x" "y" "z"
```

- ▶ You can delete an object with:

```
rm(list = c("x", "y"))
```

- ▶ You can clear your environment with:

```
rm(list = ls())
```

# Working Directory

- ▶ The working directory (WD) is the folder where R is working.
- ▶ If you read or write a file, the default location will be your WD.
- ▶ If you want to read/write from somewhere else you will have to tell R the path.
- ▶ Type the following command to find your WD:

```
getwd()
```

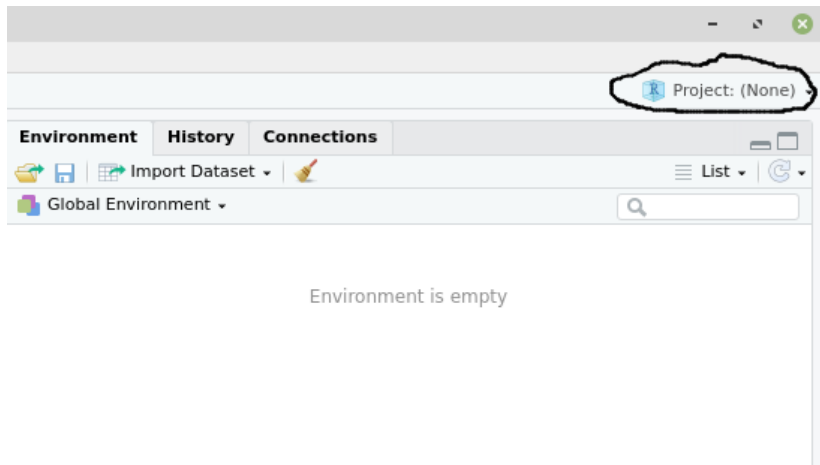
- ▶ You can change with the following command:

```
setwd("path")
```

- ▶ The path must be between quotation marks.

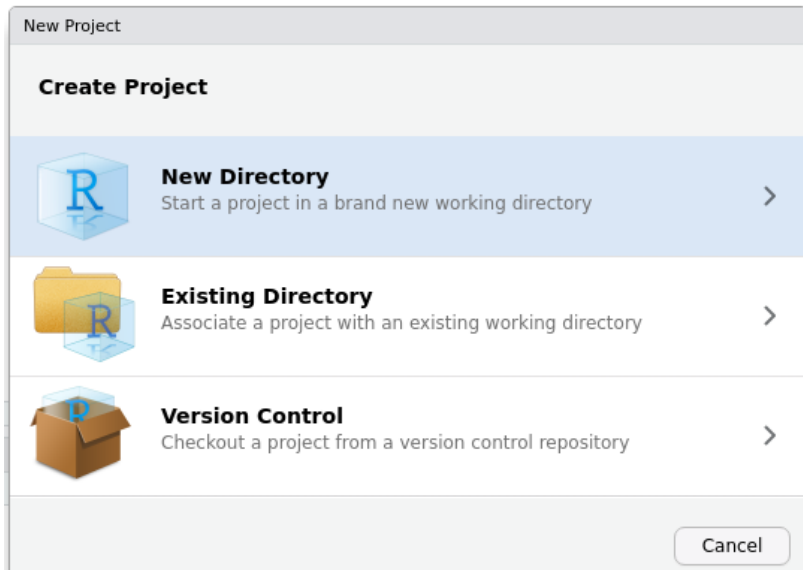
# Projects

- ▶ Projects are a very good way to keep your work organized.
- ▶ Once you setup a project it will be linked to a folder, which will be the project WD.
- ▶ You can resume your work by just opening the project and everything will be ready for you.

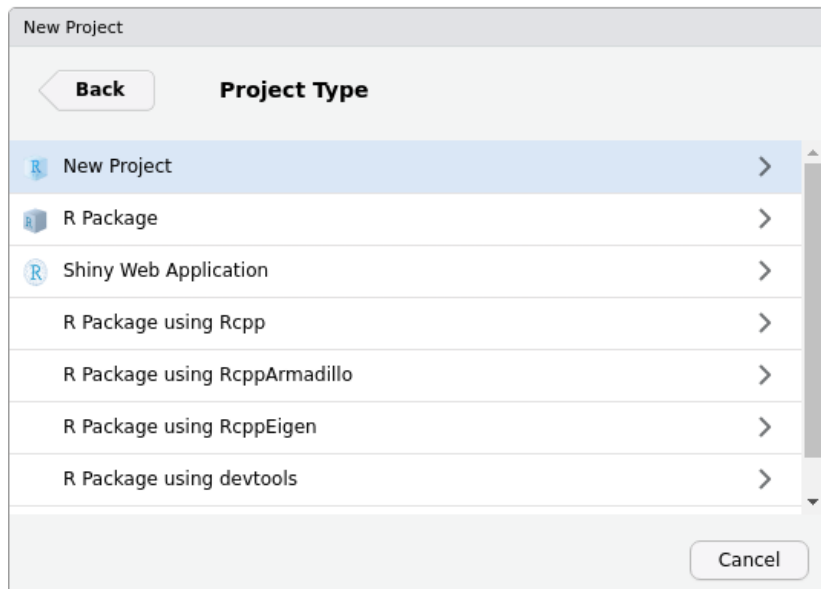


# Projects

To start a project on a new folder open file > new project and follow these steps:



# Projects




# Projects

New Project

Back

Create New Project



Directory name:  
project1


Create project as subdirectory of:  
~ Browse...

☐ Open in new session

Create Project

Cancel

# Projects





# Projects

- ▶ The project will create a file `project1.Rproj` in the selected WD. You can use it to open your project.
- ▶ You can also switch between projects by clicking on `project1` like the last slide.

# Basic Operations

```
# Addition
```

```
3 + 3
```

```
## [1] 6
```

```
# Subtraction
```

```
5 - 4
```

```
## [1] 1
```

```
# Multiplication
```

```
5 * 6
```

```
## [1] 30
```

```
# Division
```

```
10 / 5
```

```
## [1] 2
```

```
# Exponent
```

```
6 ^ 2
```

```
## [1] 36
```

# Basic Operations

- ▶ The order of priority is Exponent > Division > Multiplication > Addition = Subtraction.
- ▶ This is the order R will use in an expression like this:

```
2 + 2 * 3 + 5 / 2 ^ 2
```

```
## [1] 9.25
```

- ▶ You can change the priority using parentheses.

```
# The first addition should go first and  
# the division should go before the exponent  
(2 + 2) * 3 + (5 / 2) ^ 2
```

```
## [1] 18.25
```

## Logical/Misc Operators

Logical/Misc Operators	
==	equal to
!=	not equal equal to
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
	OR
&	AND
a%in%b	is 'a' contained in 'b'
%*%	matrix multiplication

# Logical/Misc Operators

Some examples:

```
1 == 2
```

```
## [1] FALSE
```

```
1 != 2
```

```
## [1] TRUE
```

```
3<4
```

```
## [1] TRUE
```

We will deal more with this operators latter.

# Data Types, Classes, and Structures

Up to this point, we have been working with numbers. There are actually six data types in R:

- ▶ Double: Numeric, real.
- ▶ Integer: Numeric, integer.
- ▶ Character: Name.
- ▶ logical: TRUE, FALSE
- ▶ complex:  $a + bi$  - A complex number.
- ▶ raw: a byte.

We will work only with the first four types.

# Data Types, Classes, and Structures

- ▶ You can ask R the type of the object:

```
typeof(2)
```

```
## [1] "double"
```

```
typeof(2L)
```

```
## [1] "integer"
```

```
typeof("a")
```

```
## [1] "character"
```

```
typeof(TRUE)
```

```
## [1] "logical"
```

# Data Types, Classes, and Structures

- ▶ You can also ask R if an object is of a particular type:

```
is.double(2)
```

```
## [1] TRUE
```

```
is.integer(2L)
```

```
## [1] TRUE
```

```
is.character("a")
```

```
## [1] TRUE
```

```
is.logical(TRUE)
```

```
## [1] TRUE
```

```
is.character(2)
```

```
## [1] FALSE
```



# Data Types, Classes, and Structures

- ▶ You can also ask for the class of the object.
- ▶ Class is not the same as type. For example, a matrix is a class and its type can be double, integer, character, etc. . .

```
# creates a 2x2 matrix of zeros.
```

```
A = matrix(0, 2, 2)
```

```
class(A)
```

```
## [1] "matrix"
```

```
typeof(A)
```

```
## [1] "double"
```

# Data Types, Classes, and Structures

- ▶ A particular class we will be using is called factor.
- ▶ They are used for categorical variables.
- ▶ They only accept a particular set of values called levels.

```
fact = c("a", "b", "b")  
fact = as.factor(fact)
```

```
levels(fact)
```

```
## [1] "a" "b"
```

```
class(fact)
```

```
## [1] "factor"
```

# Data Types, Classes, and Structures

- ▶ So far we mostly worked with objects of a single element.
- ▶ R have several data structures:

	Homogeneous	Heterogeneous
1-dimensional	vector	list
2-dimensional	matrix	data frame
more than 2 dimensions	array	

# Data Types, Classes, and Structures

- Vectors are created with the following command:

```
vec1 = c(2,4,6,8,10)
```

```
is.vector(vec1) # is it a vector?
```

```
## [1] TRUE
```

```
is.double(vec1) # is it type double
```

```
## [1] TRUE
```

```
is.numeric(vec1) # is it class numeric?
```

```
## [1] TRUE
```

```
is.atomic(vec1) # is it homogeneous?
```

```
## [1] TRUE
```

# Data Types, Classes, and Structures

- There are other ways to create vectors:

```
vec2 = seq(2,10,2) #seq(from,to,by)  
vec2
```

```
## [1] 2 4 6 8 10
```

```
vec3 = 1:10 #seq with by = 1  
vec3
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
length(vec3)
```

```
## [1] 10
```

# Data Types, Classes, and Structures

## ► Vector operations

```
vec4 = 1:5
```

```
vec5 = 6:10
```

```
vec4 + vec5
```

```
## [1] 7 9 11 13 15
```

```
vec4 - vec5
```

```
## [1] -5 -5 -5 -5 -5
```

```
vec4 * vec5 #inner product
```

```
## [1] 6 14 24 36 50
```

# Data Types, Classes, and Structures

- ▶ We can transpose vectors with the `t` function.

```
t(vec4)
```

```
##      [,1] [,2] [,3] [,4] [,5]  
## [1,]    1    2    3    4    5
```

```
vec4 %*% t(vec5) #5X1 times 1X5
```

```
##      [,1] [,2] [,3] [,4] [,5]  
## [1,]     6     7     8     9    10  
## [2,]    12    14    16    18    20  
## [3,]    18    21    24    27    30  
## [4,]    24    28    32    36    40  
## [5,]    30    35    40    45    50
```

```
t(vec4) %*% vec5 #1X5 times 5X1
```

```
##      [,1]  
## [1,]   130
```

# Data Types, Classes, and Structures

We can also do logical operations with vectors:

```
vec4 == vec5
```

```
## [1] FALSE FALSE FALSE FALSE FALSE
```

```
1 %in% vec4 #ask if 1 is in vec4
```

```
## [1] TRUE
```

```
vec5 > vec4
```

```
## [1] TRUE TRUE TRUE TRUE TRUE
```



# Data Types, Classes, and Structures

```
A = matrix(c(1,2,3,4,5,6),nrow=2,byrow=TRUE)
```

```
A
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    3
## [2,]    4    5    6
```

```
is.atomic(A)
```

```
## [1] TRUE
```

```
is.matrix(A)
```

```
## [1] TRUE
```

```
B = matrix(c(1,6,3,5,7,2),nrow=3)
```

```
B
```

```
##      [,1] [,2]
## [1,]    1    5
## [2,]    6    7
## [3,]    3    2
```

# Data Types, Classes, and Structures

- Here are some other matrix operations:

```
dim(A) # Dimension
```

```
## [1] 2 3
```

```
nrow(A) # Number of rows
```

```
## [1] 2
```

```
ncol(A) # Number of columns
```

```
## [1] 3
```

```
C = A%*%B # Multiplication (vectors also work)
```

```
C
```

```
##      [,1] [,2]
```

```
## [1,]    22    25
```

```
## [2,]    52    67
```

```
A*A # Inner product
```

```
##      [,1] [,2] [,3]
```

```
## [1,]     1     4     9
```

```
## [2,]    16    25    36
```

# Data Types, Classes, and Structures

► You can also try:

```
det(C) # Determinant
```

```
## [1] 174
```

```
eigen(C) # Eigenvalues/ vectors
```

```
## eigen() decomposition
```

```
## $values
```

```
## [1] 87 2
```

```
##
```

```
## $vectors
```

```
##           [,1]      [,2]
```

```
## [1,] -0.3589791 -0.7808688
```

```
## [2,] -0.9333456  0.6246950
```

```
solve(C) # Inverse
```

```
##           [,1]      [,2]
```

```
## [1,]  0.3850575 -0.1436782
```

```
## [2,] -0.2988506  0.1264368
```

# Other Built-in Functions

- R has several built-in functions:

```
mean(vec1)
```

```
## [1] 6
```

```
sd(vec1) # standard deviation
```

```
## [1] 3.162278
```

```
sum(vec1)
```

```
## [1] 30
```

```
prod(vec1)
```

```
## [1] 3840
```

```
log(vec1) #ln
```

```
## [1] 0.6931472 1.3862944 1.7917595 2.0794415 2.3025851
```

```
exp(vec1) # Exponential
```

```
## [1] 7.389056 54.598150 403.428793 2980.957987 22026.465795
```

```
summary(log(vec1))
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
## 0.6931  1.3863  1.7918  1.6506  2.0794  2.3026
```

## Indexing

- ▶ You can access a particular element of a vector/matrix using indexes:

```
vec1[2] # second element of vec1
```

```
## [1] 4
```

```
vec1[1:3] # elements 1 2 and 3 of vec1
```

```
## [1] 2 4 6
```

```
C[2,1] # row 2 column 1 of matrix C
```

```
## [1] 52
```

```
which(vec1 > 3) #position of elements in vec1 bigger than 3
```

```
## [1] 2 3 4 5
```

```
vec1[which(vec1 > 3)] # elements in vec1 bigger than 3
```

```
## [1] 4 6 8 10
```

## Creating and combining strings

The command **paste** is used to combine strings

```
names<-paste("samp",1:4,sep="")  
names
```

```
## [1] "samp1" "samp2" "samp3" "samp4"
```

```
namesWithSp<-paste("samp",1:4,sep=" ")  
namesWithSp
```

```
## [1] "samp 1" "samp 2" "samp 3" "samp 4"
```

```
namesByMach<-paste("samp",1:4,sep="Mach")  
namesByMach
```

```
## [1] "sampMach1" "sampMach2" "sampMach3" "sampMach4"
```

# Getting Help

- ▶ You can see how a function works with `?functionName`.
- ▶ This will give you access to the function documentation, which is standardized for all R functions.
- ▶ If you get an error or warning message, take a deep breath, read the error or warning, and try to figure out your error.
- ▶ If all else fails, Google the error.

# Working with data

- ▶ The most usual way to read data in R is from csv files with **read.csv**.
- ▶ However, R has tools to read many types of data like xlsx, dta, etc.
- ▶ R also has its own way of storing data (.rda and .RData files).
- ▶ This is how you can read a csv file:

```
data = read.csv("cars.csv", row.names = 1)
head(data) # shows the first 6 entries.
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

- ▶ The **rownames = 1** is to set the first column in the csv file as the names of the rows in the data. In this case we have the names of the cars.
- ▶ You can also write csv files.

```
write.csv(data, file = "path/name.csv")
```



## Working with data

- ▶ The R data extensions (.rda and .RData files) are very useful if you are dealing with large datasets.
- ▶ However, they will only work in R (other software may be compatible with adjustments).
- ▶ The advantage of these tools is that they compile the data in to a binary file, which is much faster to load and save.
- ▶ .rda is for a single object:

```
save(data, file = "path/data.rda")
```

- ▶ You can also save your entire workspace with .RData:

```
save.image("path/name.RData")
```

# Working with data

- ▶ The **summary** function gives descriptive information on the data. The statistics will be calculated for each individual column.

```
summary(data)
```

- ▶ The **View** function opens the data on a sheet similar to excel, but in read only mode.

```
View(data)
```

- ▶ You can open the same sheet by clicking at the data object in the environment (top right part of RStudio).

# Data Frames

- ▶ When we loaded the data, a new class of object was introduced:

```
class(data)
```

```
## [1] "data.frame"
```

- ▶ Data Frames are a special class made to deal with data. Each column is a variable and each row is an observation.
- ▶ It falls in the heterogeneous classes. Each column can have a different type.

# Data Frames

- ▶ Variable names play an important role in Data Frames. You can get a particular variable with the `$` element.

```
head(data$cyl)
```

```
## [1] 6 6 4 6 8 6
```

```
mean(data$cyl)
```

```
## [1] 6.1875
```

- ▶ You can also use indexes like in a vector/matrix:

```
data[1,2]
```

```
## [1] 6
```

```
data$disp[3]
```

```
## [1] 108
```

# Lists

- ▶ Lists are the most heterogeneous data in R. They can store anything.
- ▶ For example, you can have a list with a matrix, a vector, a string and even a function.

```
l1 = list(matrix = C, vector = vec1, string = "hello")
```

- ▶ Note that the elements in the list were named. This is a good practice that allow us to access the elements using their names.

```
l1$vector
```

```
## [1] 2 4 6 8 10
```

- ▶ You can also use numbers to index elements:

```
l1[[3]]
```

```
## [1] "hello"
```


- ▶ Data Frames can be seen as a special case of lists where each element is a vector of the same size.
- ▶ Lists may seem weird now, but later you will see that the output of many functions are lists.

# Basic Graphical Tools

- ▶ Graphical tools in R can be the topic for a whole course.
- ▶ The next slides will show just some basic examples.
- ▶ R has several built-in plot tools. However, if you want to go deep you should study a package called ggplot.
- ▶ ggplot is so wide that people wrote entire books just for it.

# Basic Graphical Tools


```
plot(data$hp,data$cyl)
```



# Basic Graphical Tools

```
hist(data$hp)
```


Histogram of data\$hp





# Basic Graphical Tools

```
plot(data$hp,data$mpg, col=data$cyl)  
legend("topleft",fill=c(4,6,8),legend=levels(as.factor(data$cyl)))
```



# Loop, If and Else

- ▶ Loops are codes for iterative processes. The most basic way to do it is using the **for** interface.
- ▶ A good example of something that can only be made in a loop is to create data that follow a random walk.
- ▶ A random walk is when we have

$$y_{t+1} = y_t + \varepsilon_{t+1}$$

where  $\varepsilon_t$  is an error term that we will sample from a normal distribution with mean 0 and variance 1.


```
N = 100
y = rep(NA, N)
y[1] = 0

set.seed(1)
for(i in 1:(N-1)){
  y[i+1] = y[i] + rnorm(1)
}
head(y)
```

```
## [1] 0.0000000 -0.6264538 -0.4428105 -1.2784391 0.3168417 0.6463495
```

## Loop, If and Else

```
plot(y,type = "l")
```



## Loop, If and Else

- ▶ You can also write loops that will only stop when a criterion is met with the **while** interface.
- ▶ In the example below we will keep adding a random number between 0 and 1 to  $x$  until  $x$  becomes bigger or equal to 100.
- ▶ Then we will evaluate the variable  $i$ , which tells us how many iterations were needed.

```
x = 0
i = 0
set.seed(1)
while(x < 100){
  x = x + runif(1)
  i = i+1
}
i
```

```
## [1] 194
```

# Loop, If and Else

- ▶ When one iteration depends on the previous it is hard to escape from the **for** and **while** loops.
- ▶ However, when the iterations are independent, we can use a set of functions from the **apply** family.
- ▶ For example, suppose we want to calculate the standard deviation of all columns in the car dataset;

```
# apply(data, dim, function, ... extra arguments)  
apply(data, 2, sd)
```

##	mpg	cyl	disp	hp	drat	wt
##	6.0269481	1.7859216	123.9386938	68.5628685	0.5346787	0.9784574
##	qsec	vs	am	gear	carb	
##	1.7869432	0.5040161	0.4989909	0.7378041	1.6152000	

# Loop, If and Else

- ▶ If we are dealing with lists we can use the **lapply** function:

```
# A list with the tree matrices we created  
matlist = list(A = A, B = B, C = C)  
# This will get the first line of all matrices  
lapply(matlist, function(x) x[1,] )
```

```
## $A  
## [1] 1 2 3  
##  
## $B  
## [1] 1 5  
##  
## $C  
## [1] 22 25
```

- ▶ The function **sapply** is for the cases where the input is a list (or vector) and the output is a single element:

```
sapply(matlist, sum)
```

```
##   A   B   C  
## 21  24 166
```

- ▶ **sapply** also works for vectors.

## Loop, If and Else

- ▶ **if** statements are used when we want R to do something once a certain condition is met.
- ▶ The syntax is **if (condition) {what to do}**
- ▶ We can use **if** to do the something we did in the **while** code, but with a for:

```
x = 0
set.seed(1)
for(i in 1:100000){
  x = x + runif(1)
  if(x>100){
    break
  }
}
i
```

```
## [1] 194
```

## Loop, If and Else

- ▶ It is possible to include an **else** statement with the **if**.
- ▶ In this case, if the condition is not met, R will run the code inside the **else**.

```
x = 1
if(x>2){
  z = x
}else{
  z = 0
}
z
```

```
## [1] 0
```



# Loop, If and Else

- ▶ Finally, in certain situations we can simplify the code with the **ifelse** function.

```
x = 1:10  
# ifelse(condition, if true, if false)  
ifelse(x>5 , 5, x)
```

```
## [1] 1 2 3 4 5 5 5 5 5 5
```

# User Defined Functions (UDF)

A UDF is a function created by the user. It follows the structure

**functionName = function(arguments){calculation}**

For Example:

```
subtract<-function(x,y){  
  result = x-y  
  return(result)  
}
```

- ▶ A function always stops when it reaches the **return** command.

# User Defined Functions (UDF)

► Now we can use our function:

```
subtract(10,5)
```

```
## [1] 5
```

```
subtract(x = 20, y = 10)
```

```
## [1] 10
```

# User Defined Functions (UDF)

- ▶ Now let's write a more interesting function.
- ▶ This function will create data from autoregressive models like:

$$y_{t+1} = \rho y_t + \varepsilon_{t+1}$$

- ▶ where  $0 \leq \rho \leq 1$


```
AR = function(N, rho){  
  y = rep(NA, N)  
  y[1] = 0  
  for(i in 1:(N-1)){  
    y[i+1] = rho*y[i] + rnorm(1)  
  }  
  return(y)  
}
```

- ▶ We can use our UDF in an sapply loop with several values of  $\rho$

```
rho_vector = c(0, 0.5, 0.95, 0.99, 1)  
result = sapply(rho_vector, function(x) AR(200,x))
```

# User Defined Functions (UDF)

```
matplot(result, type = "l")
```



# Generation of Random Numbers

- ▶ Several codes so far had a **set.seed** before the functions we were running.
- ▶ This seed is used when we are dealing with random numbers. It allows some other user to replicate the exact same experiment.
- ▶ For example, we are going to sample 3 numbers from a vector with a seed set to 1.
- ▶ This is a random experiment, but if you use the same seed in your computer you should get the same result.

```
vec6 = 1:10  
# sample(x, size, replace = FALSE)  
set.seed(1)  
sample(vec6, 3, replace = FALSE)
```

```
## [1] 9 4 7
```

# Generation of Random Numbers

- R has build in functions to generate data from several distributions like normal, uniform, Student t, etc.

```
# 10 numbers from a normal distribution with mean 0 and sd 1
```

```
rmnorm(10,0,1)
```

```
## [1] 1.329799263 1.272429321 0.414641434 -1.539950042 -0.928567035
```

```
## [6] -0.294720447 -0.005767173 2.404653389 0.763593461 -0.799009249
```

```
# 10 numbers from a uniform distribution with min 0 max 1
```

```
runif(10,-1,1)
```

```
## [1] -0.74888981 -0.46555866 -0.22777181 -0.97321933 -0.23522409 0.73938169
```

```
## [7] -0.31930201 -0.03583977 0.19913165 -0.01291739
```

```
# 10 numbers from a t distribution 5 degrees of freedom
```

```
rt(10,5)
```

```
## [1] -0.86409523 -1.50276529 0.85199410 -1.82436807 -0.06641194 -1.41288461
```

```
## [7] -0.32612422 0.44183505 0.95343054 -0.33398807
```

# Generation of Random Numbers

- ▶ Note that all functions that we just used started with a **r**, which comes from random.
- ▶ If you use a **d** you will calculate the density, **p** for the distribution and **q** for the quantile.



# Packages

- ▶ Registered packages are stored in the Comprehensive R Archive Network (CRAN).
- ▶ These packages can be installed with `install.packages("pkg name")`
- ▶ And they can be loaded with **`library(pkg name)`**.
- ▶ Once you load the package you have access to all its functions, data and documentation.

```
install.packages("glmnet")  
library(glmnet)  
?glmnet
```

# The Tidyverse

- ▶ The Tidyverse is a set of dozens of packages, all compatible with each other, made for data treatment and analysis.
- ▶ It is the state of the art for data treatment. You can install and load it with:

```
install.packages("tidyverse")  
library(tidyverse)
```

# The Tidyverse

- ▶ One of the most important features of Tidyverse is the pipe operator `%>%`.
- ▶ Consider the code:

```
log(exp(sin(2^2)))
```

```
## [1] -0.7568025
```

- ▶ Once we keep using functions inside functions the code may look very confusing and it is very easy to get lost in the parentheses.
- ▶ Also, it is important that your code can be read and understood by other humans, and not only machines.
- ▶ The same results can be obtained with the pipe operator:

```
2^2 %>% sin() %>% exp() %>% log()
```

```
## [1] -0.7568025
```

# The Tidyverse

- ▶ The pipe operator can be used in the same way with data frames. It is very fast and you can perform a lot of operations in a single step.
- ▶ For example, we can select columns:

```
df1 = data %>% select(cyl,mpg,hp)
head(df1)
```

##	cyl	mpg	hp
## Mazda RX4	6	21.0	110
## Mazda RX4 Wag	6	21.0	110
## Datsun 710	4	22.8	93
## Hornet 4 Drive	6	21.4	110
## Hornet Sportabout	8	18.7	175
## Valiant	6	18.1	105

# The Tidyverse

- ▶ We can filter the data with some criterion:

```
df2 = df1 %>% filter(cyl == 4)
head(df2)
```

```
##   cyl  mpg  hp
## 1   4 22.8  93
## 2   4 24.4  62
## 3   4 22.8  95
## 4   4 32.4  66
## 5   4 30.4  52
## 6   4 33.9  65
```

# The Tidyverse

- It is also possible to create new variables:

```
df3 = df2 %>% mutate(newvar = hp/mpg)
head(df3)
```

```
##   cyl  mpg  hp  newvar
## 1    4 22.8 93 4.078947
## 2    4 24.4 62 2.540984
## 3    4 22.8 95 4.166667
## 4    4 32.4 66 2.037037
## 5    4 30.4 52 1.710526
## 6    4 33.9 65 1.917404
```

# Tidyverse

- The **group\_by** and the **summarise** functions together are used to group the observations by some variable.

```
df4 = df1 %>% group_by(cyl) %>%  
  summarise(hp = mean(hp), mpg = mean(mpg))  
head(df4)
```

```
## # A tibble: 3 x 3  
##   cyl    hp  mpg  
##   <int> <dbl> <dbl>  
## 1     4  82.6  26.7  
## 2     6 122.   19.7  
## 3     8 209.   15.1
```

# Tidyverse

- ▶ Finally, we can combine as many operations as we want in one chain of pipe codes.

```
df5 = data %>% select(mpg, hp, cyl) %>%  
  filter(cyl > 4) %>%  
  group_by(cyl) %>%  
  summarise(hp = mean(hp), mpg = mean(mpg)) %>%  
  mutate(newvar = hp/mpg)
```

df5

```
## # A tibble: 2 x 4  
##   cyl    hp    mpg newvar  
##   <int> <dbl> <dbl> <dbl>  
## 1     6  122.  19.7   6.19  
## 2     8  209.  15.1  13.9
```